

Data Programming with R

Isabella Gollini

Lecture 6 - R programming structures

R programming structures

- Programming in R
 - loops - if-else
 - Boolean operators
- Writing functions
 - default values
 - `return()` and `invisible()`
 - functions are objects
 - viewing and editing flexible arguments
- Environments and scope
 - `ls` and `ls.str`
 - global assignment
- Clever coding
 - recursion
 - replacement functions
 - writing operators
 - anonymous functions

Programming in R

- Without mentioning it, we've learned quite a bit about the way R works:
 - In R everything is an object (even a function).
 - Functions have arguments given in brackets.
 - New statements are given a new line but can alternatively be separated by a semi-colon.
 - There's no need to declare (most types of) variables in R.
- In this lecture we will go into the structure of language in a bit more detail and cover a few more important concepts.

Loops

Looping in R is slow but sometimes necessary. The simplest loop is a `for` loop:

```
for(i in 1:2) print(i)
```

```
## [1] 1
## [1] 2
```

This says: set `i` to be 1 and then carry out instructions, then set `i` to be 2 and carry out instructions, etc

We can make this richer by expressing the looping vector to be an object:

```
x <- c(2, 4, 6)
for(i in x) { # for loop over the elements of x
  print(i) # print the current value of i
  print(i^2) # print the current value of i^2
}
```

```
## [1] 2
## [1] 4
## [1] 4
```

```
## [1] 16
## [1] 6
## [1] 36
```

`i` is set to be `x[1]`, `x[2]`, etc and the curly braces encapsulate a block of commands to be carried out.

Note that it is common to indent text inside curly brackets to make it easier to read.

while and repeat loops

while loops keep repeating until a certain condition is met.

repeat loops keep repeating until a `break` command is found.

The `break` command can also be used in `for` and `while` loops.

```
i <- 2 # initialise i = 2
while(i < 10){ # stop if i >= 10
  print(i)
  print(i^2)
  i <- i + 2
}
```

```
## [1] 2
## [1] 4
## [1] 4
## [1] 16
## [1] 6
## [1] 36
## [1] 8
## [1] 64
```

```
i <- 2 # initialise i = 2
repeat{ # repeat the code in curly brackets
  print(i)
  print(i^2)
  i <- i + 2
  if(i > 8) break # exit the loop if i > 8
}
```

```
## [1] 2
## [1] 4
## [1] 4
## [1] 16
## [1] 6
## [1] 36
## [1] 8
## [1] 64
```

Another useful command is `next` which tells a loop to skip to the next iteration without conducting any of the later commands.

Looping over non-vectors

Almost all loops take place over a vector. If you want to loop over something else (e.g. a list of objects or a matrix) you need to do something smarter.

Example: looping over matrix objects

```

u <- matrix(1:6, 3, 2)
v <- matrix(7:12, 3, 2)
for(i in c('u', 'v')) {
  print(i) # print the value of i
  z <- get(i) # get the object with the name given by the current value of i
  print(colMeans(z)) # print the values of the column means
}

## [1] "u"
## [1] 2 5
## [1] "v"
## [1] 8 11

```

The (very powerful) `get` function here turns the character vector (`u`, `v`) into objects from their names.

if-else statements

We have already seen simple if statements used in commands. A useful extension is if-else:

```

r <- 3
if(r == 4) {
  x <- 2
} else { # if( r != 4)
  x <- 3
  y <- 4
}

```

- If you have multiple commands to carry out inside an if statement you should always use curly brackets.
- Remember that the value inside the `if` statement is a logical condition (i.e. must evaluate to `TRUE/FALSE`).

Arithmetic and Boolean operators

We are already used to writing e.g. `x + y`, `x - y`, `x * y`, `x / y`; and `x^y` and `x %% y`

Another useful command is `x %/% y` which is for integer division, e.g. `12 %/% 7 = 1` because 7 goes into 12 once.

Boolean operations include `x == y` (equality), `x > y`, `x < y`, `x >= y`, `x <= y` (inequalities), `x != y` (negation).

There are also **AND** and **OR** Boolean operations for Boolean variables: `x && y`, `x || y` which produce scalar output and `x & y`, `x | y` which produce vector output.

```

a <- c(3, 2, 1)
b <- c(1, 2, 3)
# check if: (a[1] > 1 AND b[1] < 2); (a[2] > 1 AND b[2] < 2); (a[3] > 1 AND b[3] < 2)
a > 1 & b < 2

## [1] TRUE FALSE FALSE

# check only if: a[1] > 1 AND b[1] < 2
a > 1 && b < 2

## [1] TRUE

# check if: (a[1] > 1 OR b[1] < 2); (a[2] > 1 OR b[2] < 2); (a[3] > 1 OR b[3] < 2)
a > 1 | b < 2

## [1] TRUE TRUE FALSE

```

```
# check only if: a[1] > 1 OR b[1] < 2
a > 1 || b < 2
```

```
## [1] TRUE
```

The difference is that `x && y` and `x || y` look at only the first elements of `x` and `y`. It's usually safer to use `x & y` or `x | y`.

More on Boolean operations

`TRUE` and `FALSE` can be shortened to `T` and `F` provided you have not used the variables `T` and `F` elsewhere in your code (try not to!).

`TRUE` maps to (i.e. is stored by R as) the value 1 and `FALSE` as 0.

```
TRUE + 2
```

```
## [1] 3
```

```
TRUE * 5
```

```
## [1] 5
```

```
(1 < 2) == 1 # same as: TRUE == 1
```

```
## [1] TRUE
```

```
FALSE * 5
```

```
## [1] 0
```

R has a specific mode for Boolean variables: `logical`. There are associated `as.logical` and `is.logical` functions available.

Writing functions

- default values
- `return()` and `invisible()`
- functions are objects
- viewing and editing flexible arguments

Writing functions

Here is a simple function with 3 arguments:

```
fun1 <- function(x, printx = TRUE, printy = TRUE) {
  y <- x^2
  if(printx) print(x) # if(printx) is the same as if(printx == TRUE)
  if(printy) print(y) # if(printy) is the same as if(printy == TRUE)
  return(y)
}
```

We can run the function with, e.g.:

```
answer <- fun1(5) # by default it prints x and y
```

```
## [1] 5
```

```
## [1] 25
```

```
answer # show the value returned by fun1
```

```
## [1] 25
```

Default values for arguments

The first line of the function specifies that `printx` and `printy` have default arguments (both set to `TRUE`). The argument `x` has no default and must be supplied.

If we want to turn the `printx` and `printy` arguments off we can specify them:

```
answer <- fun1(5, printx = FALSE, printy = FALSE) # does not print x and y
answer # show the value returned by fun1
```

```
## [1] 25
```

We can give the arguments in any order provided we specify the argument names:

```
answer <- fun1(printy = FALSE, x = 7) # does not print y, and the value of x is 7
```

```
## [1] 7
```

```
answer
```

```
## [1] 49
```

However if we don't give the argument names we must specify them in order:

```
answer <- fun1(5, FALSE) # this is the same as fun1(x = 5, printx = FALSE, printy = TRUE)
```

```
## [1] 25
```

```
answer
```

```
## [1] 25
```

return or not?

At the end of the function I have used `return()` to explicitly tell R that this is the value we want to give back. Thus if we write:

```
answer <- fun1(5)
```

```
## [1] 5
```

```
## [1] 25
```

then `answer` will store the value(s) specified by `return`.

As we are only returning a single value, we can ignore `return` and just have `y` as the last line.

If we want to return multiple values we can return them as a `list`:

```
fun2 <- function(x, printx = TRUE, printy = TRUE) {
  y <- x^2
  if(printx) print(x)
  if(printy) print(y)
  return(list(input = x, output = y))
}
```

Although not compulsory it's usually a good idea to put `return` at the bottom of the function.

invisible

By default the function will print out the answer every time:

```
fun1(5, FALSE, FALSE)
```

```
## [1] 25
```

If we use `invisible` instead of `return` we can remove this behaviour:

```
fun3 <- function(x, printy = TRUE) {
  y <- x^2
  if(printy) print(y)
  invisible(y) # do not print the value returned by the function if that's
               # not assigned to an object
}
fun3(5, FALSE)
answer <- fun3(5, FALSE)
answer # the value is returned since is saved in 'answer'

## [1] 25
```

The function function

The R function `function` has essentially two arguments, the argument list and the body

These can be accessed via the functions `formals` and `body`

Consider:

```
fun4 <- function(x = 7){
  return(x^2)
}
formals(fun4) # get the arguments of the function fun4
```

```
## $x
## [1] 7
```

```
body(fun4) # get the body of the function fun4
```

```
## {
##   return(x^2)
## }
```

Viewing and editing R functions

- When you type in the name of a function without brackets R will print out the arguments and body of the function.
- This works well for functions that you have written and for some R functions (see, e.g. `abline`, `read.table`) but not for others which are programmed in other languages (e.g. `mean`, `sum`).
- If a function is very long you can use `page(function)` to view it more neatly, or `edit(function)` to edit it.

If you're going to edit an existing R function it is often best to rename it first:

```
abline2 <- abline
page(abline2) # show the function in a "R page"
edit(abline2) # get access to the text editor to edit the function 'abline'
```

- Notice that function `abline` uses `invisible` instead of `return`

Flexible arguments in functions

- Many of the default R functions (e.g. `plot`) contain an ellipsis `...` at the end of the argument list

This allows you to pass in other extra named arguments to functions that may be called as part of the main function.

```
fun5 <- function(x,...) {
  return(mean(x,...))
}
```

Here mean has the extra arguments `na.rm` and `trim`, so:

```
fun5(c(1:10))
```

```
## [1] 5.5
```

```
fun5(c(1:10, NA))
```

```
## [1] NA
```

```
fun5(c(1:10, NA), na.rm = TRUE)
```

```
## [1] 5.5
```

So I can give `fun5` the argument `na.rm = TRUE` despite it not being detailed in the original function arguments.

Environments and scope



- `ls` and `ls.str`
- global assignment
- Those of you without a programming background might prefer to watch/read this section twice. See the screencast for an extra walkthrough.

Consider the following example:

```
w <- 12
f <- function(y) { # create the function f
  d <- 8
  h <- function(){ # create the function h inside the function f
    return(d * (w + y)) # d is passed by the environment of f
                        # w is passed by the global environment
                        # y is passed as an argument of f
  }
  return(h()) # function f returns the value of h()
}
f(2)
```

- Without running this code, what do you think the value of `f(2)` will be?
- When you are at the standard R prompt you are in the 'Global Environment'. Any objects you create here are available to every function you might create.
- Once you are inside a function, however, any objects you create are only available to that function. They are created in a separate environment.
- Functions can be created inside other functions creating nested environments where only certain objects are available.
- The `environment` function will tell you which environment you are in when you run it.
- The `ls` (and `ls.str`) functions will list all the objects in the current environment

environment, `ls` and `ls.str`

- Returning to our example function:

```
environment() # show the current environment
```

```
## <environment: R_GlobalEnv>
```

```
ls() # give the list the objects in the current environment
```

```
## [1] "f" "w"
```

```
ls.str() # give the list the objects in the current environment and their structure
```

```
## f : function (y)
```

```
## w : num 12
```

- The global environment (i.e. available to everything) is known as `.GlobalEnv`, though confusingly the `environment` function calls it `R_GlobalEnv`.
- Note that `h` and `d` don't exist as they are only created inside the function so are not in the global environment.

Environments inside functions

- Now suppose we change our function slightly to print the environment whilst we're in the function

```
f2 <- function(y) {  
  d <- 8  
  h <- function(){  
    return(d * (w + y))  
  }  
  print(environment())  
  print(ls.str())  
  return(h())  
}
```

```
f2(2)
```

```
## <environment: 0x7fc6c6cf2920>
```

```
## d : num 8
```

```
## h : function ()
```

```
## y : num 2
```

```
## [1] 112
```

Getting it wrong

Suppose instead we gave the functions separately:

```
f <- function(y) {  
  d <- 8  
  return(h())  
}
```

```
h <- function() {  
  return(d * (w + y))  
}
```

```
f(2)
```

```
## Error in h(): object 'd' not found
```



```
f <- function(y) {
  d <- 8
  return(h())
}
# change the order in the multiplication in h
h <- function() {
  return((w + y) * d)
}

f(2)
```

```
## Error in h(): object 'y' not found
```

- These errors are returned because the objects `d` and `y` do not exist in the environment for `h`.
- The fix is to give `h` some named arguments - see the screencast (mentioned in slide 20) for details.
- Notice that R stops the function when the first error it finds.

More on `ls`.

- So far we have called `ls` without any arguments to give the list of objects in the current environment.

By adding the `envir` argument, we can list all the objects in environments above that of where `envir` is called.

```
f <- function(y) {
  d <- 8
  return(h(d, y))
}
h <- function(dd, yy) {
  print(ls()) # list the arguments in the current environment
  print(ls(envir = parent.frame(n = 1))) # list the arguments in the environment above
  return(dd * (w + yy))
}

# When h is called from f ls(envir = parent.frame(n = 1)) show the arguments in f
f(2)
```

```
## [1] "dd" "yy"
## [1] "d" "y"
## [1] 112
```

```
# When h is used directly ls(envir = parent.frame(n = 1)) show the arguments
# in the global environment
h(8, 2)
```

```
## [1] "dd" "yy"
## [1] "f" "f2" "h" "w"
## [1] 112
```

Objects can exist only in functions

If you change an object inside a function the object will remain unchanged once the function has finished its execution

```
d <- 5
f <- function() {
  d <- 8
```

```
    return(d)
  }
  f()
```

```
## [1] 8
```

```
d
```

```
## [1] 5
```

- This behaviour occurs because `d` lives only inside the function `f`
- However, there is an exception to this rule...

Global assignment

The `<-` and `assign` functions allow you to change an object inside a function and have it persist in different environments:

```
d <- 5
f <- function() {
  d <- 8
  return(d)
}
f()
```

```
## [1] 8
```

```
d
```

```
## [1] 8
```

```
d <- 5
f <- function() {
  d <- 8
  assign('d', d, pos = .GlobalEnv)
  return(d)
}
f()
```

```
## [1] 8
```

```
d
```

```
## [1] 8
```

- The extra advantage of using `assign` is that the first argument is a character string so you can assign a vector of names to different values and then call them all back in with `get`

Clever coding

- recursion
- replacement functions
- writing operators
- anonymous functions

Recursion

- A recursive function is one that calls itself.
- This can be surprisingly useful and efficient though is often confusing for those who are not used to writing functions.
- Consider finding

$$\begin{aligned}x! &= x \times (x-1) \times (x-2) \times \dots \times 1 \\ &= x \times (x-1)!\end{aligned}$$

A simple function is:

```
fact <- function(x) {  
  if(x == 1) return(x)  
  return(x * fact(x - 1))  
}  
fact(5)
```

```
## [1] 120
```

- Note that there is a break as soon as `x` becomes 1, otherwise the function would run on forever.
- The `return` function breaks the problem down into a smaller problem and re-runs the function.

Quick sort

- Another nice recursion example is the quick sort algorithm.

Given a vector, e.g. (5, 4, 12, 13, 3, 8, 88) this algorithm compares the first element 5 with the others and produces two sub-vectors; the first with all the elements smaller than it 4, 3, the second with all the elements bigger than it 12, 13, 8, 88 and re-runs `quicksort` on these two sub-vectors.

```
quicksort <- function(x) {  
  if(length(x) <= 1) return(x)  
  pivot <- x[1] # take the first value as pivot  
  therest <- x[-1] # create a new vector with all values but the pivot  
  sv1 <- therest[therest < pivot] # subset of values smaller than the pivot  
  sv2 <- therest[therest >= pivot] # subset of values greater than or = the pivot  
  sv1 <- quicksort(sv1) # quicksort the values smaller than the pivot  
  sv2 <- quicksort(sv2) # quicksort the values greater than or equal the pivot  
  return(c(sv1, pivot, sv2))  
}  
quicksort(c(5, 4, 12, 13, 3, 8, 88))
```

```
## [1] 3 4 5 8 12 13 88
```

Replacement functions

R allows you to use replacement functions in certain circumstances:

```
names(x)
```

```
## NULL
```

```
names(x) <- c('a', 'b', 'ab')
```

```
names(x)
```

```
## [1] "a" "b" "ab"
```

- The `names(x) <- ...` line looks like it violates some of our rules - we're forcing the outcome of a function to be a set of values.
- Bizarrely `names<-` is actually a function! (Side note: this is one of those situations where `<-` works and `=` does not.)

What we're really calling is:

```
x <- 'names<-'(x, c('a', 'b', 'ab'))
```

- Whenever you see a function on the left hand side this is really a replacement function of the form above.

Writing your own operator

If you want to write your own operator, you simply need to start and end your function with `%` and enclose it in quotes:

```
"%a2b%" <- function(a, b) return(a + 2 * b)
3 %a2b% 5
```

```
## [1] 13
```

- Note that your function must take two arguments and return a single value.

Anonymous functions

Sometimes it's not even worth giving your function a name. Eg. if you're just using it and disposing of it instantly:

```
M <- matrix(1:6, nrow = 3, ncol = 2)
f <- function(x) x / max(x)
apply(M, 1, f)
```

```
##      [,1] [,2] [,3]
## [1,] 0.25  0.4  0.5
## [2,] 1.00  1.0  1.0
```

```
apply(M, 1, function(x) x / max(x))
```

```
##      [,1] [,2] [,3]
## [1,] 0.25  0.4  0.5
## [2,] 1.00  1.0  1.0
```

- In the second example we haven't had to define the function at all!

Lessons from this week

- Lots of important things this week!
- Loop through with `for`, `while`, `repeat`, but avoid them whenever you can.
- Be careful when writing functions to specify your arguments correctly and return your output in the way that you want it.
- Pay close attention to the environment in which your function is called.